


Artificial Neural Networks



Feed Forward Networks

- Applications

- Classical Examples

Multi Layer Networks

- Possible Mappings

- Backprop Algorithm

- Practical Problems

Deep Networks

- Vanishing Gradients

- Convolutional Networks

Feed Forward Networks

- Applications

- Classical Examples

Multi Layer Networks

- Possible Mappings

- Backprop Algorithm

- Practical Problems

Deep Networks

- Vanishing Gradients


- Convolutional Networks

Artificial Neural Networks (ANN)

- ▶ Inspired from the nervous system
- ▶ Parallel processing

We will focus on **one** class of ANNs:

Feed-forward Layered Networks



Applications

Operates like a general "Learning Box"!

Classification



Function Approximation



Multidimensional Mapping



Classical Examples

ALVINN

Autonomous driving




Trained to mimic the behavior of human drivers

Classical Examples

NetTalk

Speech Synthesis



Trained using a large database of spoken text

Feed Forward Networks

Applications

Classical Examples

Multi Layer Networks

Possible Mappings

Backprop Algorithm



Practical Problems

Deep Networks

Vanishing Gradients

Convolutional Networks

What is the point of having multiple layers?



A two layer network can implement **arbitrary decision surfaces**
...provided we have *enough hidden units*

How can we train a multi layer network?



Neither perceptron learning, nor the delta rule can be used


Fundamental problem:

When the network gives the wrong answer
there is no information on in which direction
the weights need to change to improve the result


Trick:

Use threshold-like, but **continuous** functions






First layer response



Second layer response



Learning strategy:

Minimize the error (E) as a function of **all** weights (\vec{w})

1. Compute the direction in weight space where the error increases the most $\text{grad}_{\vec{w}}(E)$
2. Change the weights in the opposite direction


$$w_i \leftarrow w_i - \eta \frac{\partial E}{\partial w_i}$$

Normally one can use the error from each example separately
(Stochastic Gradient Descent)

$$E = \frac{1}{2} \sum_{k \in \text{Out}} (t_k - y_k)^2$$

A common "threshold-like function" is

$$\rho(x) = \frac{1}{1 + e^{-x}}$$



The gradient can be expressed as a function of a *local generalized error* δ

$$\frac{\partial E}{\partial \mathbf{w}_{ji}} = -\delta_i \mathbf{x}_j \quad \mathbf{w}_{ji} \leftarrow \mathbf{w}_{ji} + \eta \delta_i \mathbf{x}_j$$

Output layer:

$$\delta_k = y_k \cdot (1 - y_k) \cdot (t_k - y_k)$$

Hidden layers:

$$\delta_j = h_j \cdot (1 - h_j) \cdot \sum_k \mathbf{w}_{kj} \delta_k$$

The errors δ propagate backwards through the layers

Error backpropagation (BackProp)

Things to think about when using BackProp

- ▶ Sloooooow
Normal to require thousands of iterations through the dataset
 - ▶ Gradient following
Risk of getting stuck in local minima
 - ▶ Many parameters
 - ▶ Step size η
 - ▶ Number of layers
 - ▶ Number of hidden units
 - ▶ Input and output representation
 - ▶ Initial weights
-

Feed Forward Networks

- Applications

- Classical Examples

Multi Layer Networks

- Possible Mappings


- Backprop Algorithm

- Practical Problems

Deep Networks

- Vanishing Gradients

- Convolutional Networks




Deep networks — Networks with many layers

- ▶ Error gradients become smaller from layer to layer
- ▶ Pure Backprop becomes unusable for deep networks

Deep Belief Networks

- ▶ Unsupervised learning of features
Restricted Boltzmann Machine
 - ▶ Greedy learning from the bottom, layer by layer
Optimize for ability to reconstruct previous layer
 - ▶ Supervised Backprop to finalize classifier
-

Convolutional Networks



- ▶ Alternating convolution and subsampling layers
- ▶ Weight sharing
- ▶ Trained using Backprop